

# Kubernetes Security - en

## Agenda

### 1. Starting

- [The truth about security](#)
- [The architecture of Kubernetes](#)
- [Architecture DeepDive](#)
- [Layers to protect \(Security\)](#)
- [AttackVectors](#)
- [The route from development to production to secure](#)
- [Kill Chain](#)

### 2. Getting hacked

- [Why is a cluster so rewarding to hack](#)
- [Starting with Tesla](#)

### 3. Category 1 by Layer: OS / Kernel

- [Securing the OS and the Kernel](#)
- [Kernel Hardening Checker](#)

### 4. Category 2 by Layer: Cluster

- [Securing the components](#)
- [Securing kubelet](#)
- [Least Privileges with RBAC](#)
- [Admission Controller](#)

### 5. Category 3 by Layer: Pods Container

- [The runAs Options in SecurityContext](#)
- [sysctls in pods/containers](#)
- [Overview capabilities](#)
- [Start pod without capabilities & how can we see this](#)
- [Hacking and exploration session HostPID](#)
- [Great but still alpha User Namespaces](#)

### 6. Reaction

- [The Audit Logs](#)

### 7. RBAC

- [How does RBAC work ?](#)
- [Where does RBAC play a role ?](#)
- [kubeconfig decode certificate](#)
- [kubectl check your permission - can-i](#)
- [use kubectl in pod - default service account](#)
- [create user for kubeconfig with using certificate](#)
- Components / moving parts of RBAC
- [practical exercise rbac](#)

### 8. Obey Security Policies (AdmissionControllers)

- [Admission Controller](#)
- PSA (PodSecurity Admission)
- [Exercise with PSA](#)
- [OPA Gatekeeper 01-Overview](#)
- [OPA Gatekeeper 02-Install with Helm](#)
- [OPA Gatekeeper 03-Simple Exercise](#)
- [OPA Gatekeeper 04-Example-Job-Debug](#)
- [Connaissance: Verifying images before Deployment](#)

### 9. Pod Security

- [Automount ServiceAccounts or not ?](#)
- Does every pod need to access the kubernetes api server?

### 10. Unprivileged Pods/Container

- [Which images to use ?](#)
- How can i debug non-root - Container/Pods?

### 11. The SecurityContext

- seccomp
- privileged/unprivileged
- appArmor / SELinux

### 12. Network Policies

- Understand NetworkPolicies
- [Exercise NetworkPolicies](#)

#### 13. ServiceMesh

- [Why a ServiceMesh ?](#)
- [How does a ServiceMeshs work? \(example istio\)](#)
- [istio security features](#)
- [istio-service mesh - ambient mode](#)
- [Performance comparison - baseline,sidecar,ambient](#)

#### 14. Image Security

- When to scan ?
- Image Security Scanning

#### 15. Documentation

- [Great video about attacking kubernetes - older, but some stuff is still applicable](#)
- [Straight forward hacking session of kubernetes](#)
- [github with manifests for creating bad pods](#)

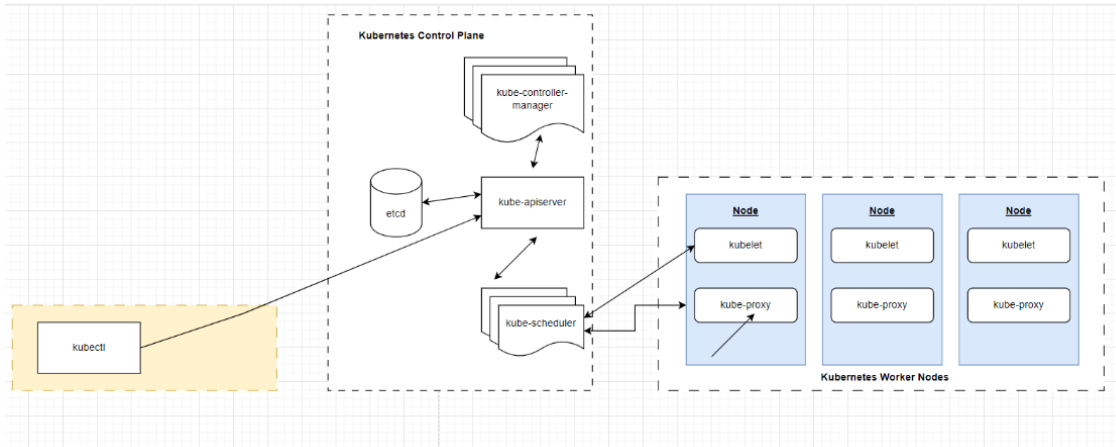
## Starting

### The truth about security

- It is an ongoing process
- Kubernetes is not safe by default

### The architecture of Kubernetes

#### Overview



## Components

### Master (Control Plane)

#### Jobs

- The master coordinates the cluster
- The master coordinates the activities in the cluster
  - scheduling of applications
  - to take charge of the desired state of application
  - scaling of applications
  - rollout of new updates

### Components of the Master

#### ETCD

- Persistent Storage (like a database), stores configuration and status of the cluster

#### KUBE-CONTROLLER-MANAGER

- In charge of making sure the desired state is achieved (done through endless loops)
- Communicates with the cluster through the kubernetes-api (kube-api-server)

#### KUBE-API-SERVER

- provides api-frontend for administration (no gui)
- Exposes an HTTP API (users, parts of the cluster and external components communicate with it)
- REST API

#### KUBE-SCHEDULER

- assigns Pods to Nodes.
- scheduler determines which Nodes are valid placements for each Pod in the scheduling queue ( according to constraints and available resources )
- The scheduler then ranks each valid Node and binds the Pod to a suitable Node.
- Reference implementation (other schedulers can be used)

### Nodes

- Nodes (Knoten) sind die Arbeiter (Maschinen), die Anwendungen ausführen
- Ref: <https://kubernetes.io/de/docs/concepts/architecture/nodes/>

### Pod/Pods

- pods are the smallest unit you can roll out on the cluster
- a pod (basically another word for group) is a group of 1 or more containers
  - mutually used storage and network resources (all containers in the same pod can be reached with localhost)
  - They are always on the same (virtual server)

### Control Plane Node (former: master) - components

### Node (Minion) - components

#### General

- On the nodes we will rollout the applications

#### kubelet

```
Node Agent that runs on every node (worker)
its job is to download images and start containers
```

#### kube-proxy

- Runs on all of the nodes (DaemonSet)
- Is in charge of setting up the network rules in iptables for the network services
- Kube-proxy is in charge of the network communication inside of the cluster and to the outside

#### ref:

- <https://www.redhat.com/en/topics/containers/kubernetes-architecture>

#### Architecture DeepDive

- <https://github.com/jmetzger/training-kubernetes-advanced/assets/1933318/1ca0d174-f354-43b2-81cc-67af8498b56c>

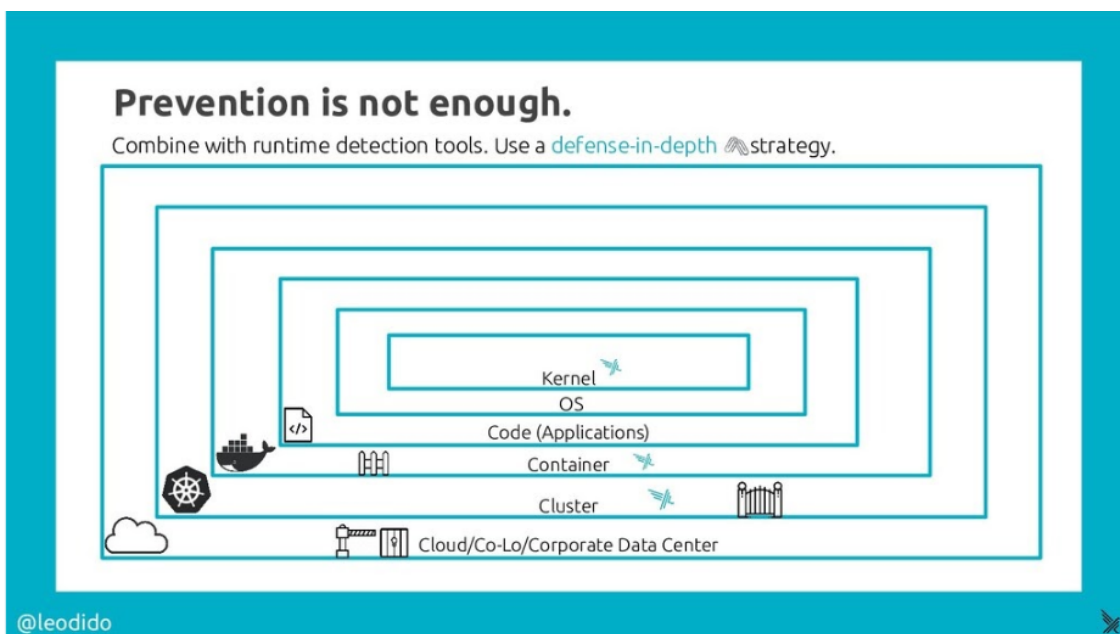
#### Layers to protect (Security)

##### Based on the 4-C - Model

- Cloud
- Cluster
- Container
- Code

##### But let us put it a bit further:

- OS
- Kernel



Credits: @Leodido

#### AttackVectors

##### What 3 types of attack vectors are there ?

1. External Attackers
2. Malicious containers
3. Compromised or malicious users

##### External Attackers

- You can have threat actors who have no access to the cluster but are able to reach the application running on it.

##### Malicious containers

- If a threat actor manages to breach a single container, they will attempt to increase their access and take over the entire cluster.

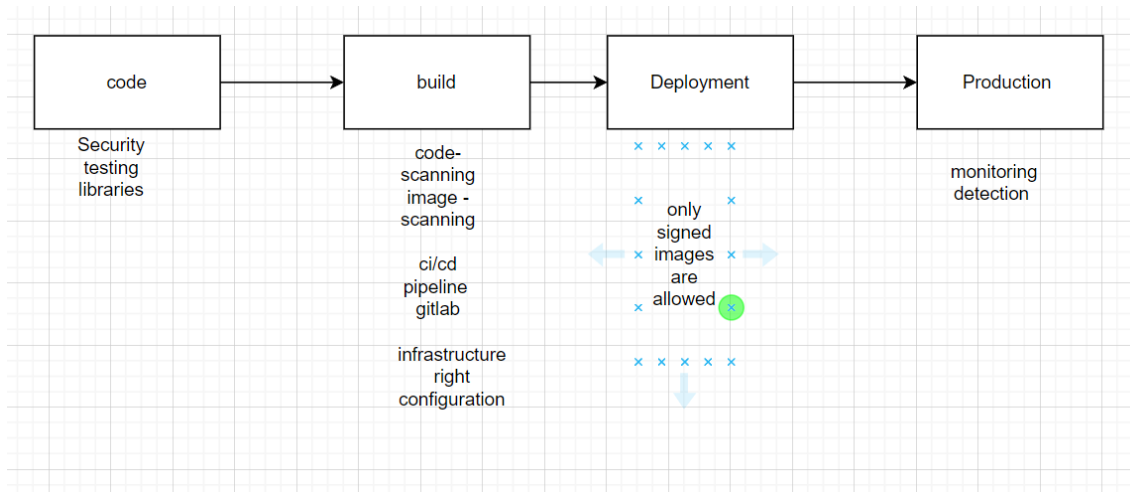
##### Compromised or malicious users

- When you're dealing with compromised accounts or malicious users, an attacker with stolen yet valid credentials will execute commands against network access and the Kubernetes API.
- Mitigation: Least Principle Policy

##### Reference:

- <https://www.cncf.io/blog/2021/11/08/kubernetes-main-attack-vectors-tree-an-explainer-guide/>

## The route from development to production to secure



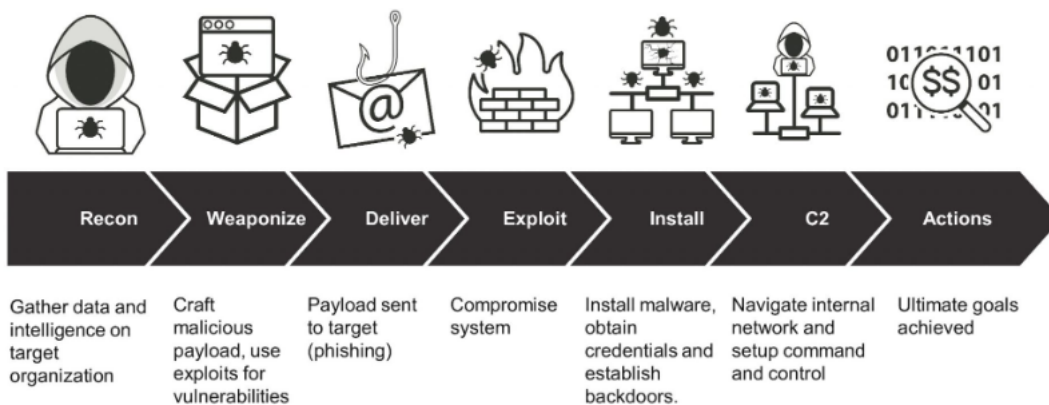
## Kill Chain

### Steps

1. Reconnaissance
2. Weaponization (Trojaner)
3. Delivery (wie liefern wie ihn aus ?)
4. Exploit (Sicherheitslücke ausnutzen)
5. Installation (phpshell)
6. Command & Control
7. Action/Objectives (mein Ziel)

Or better: graphical

## Traditional Kill Chain Model



(Source: techtag.de)

## Getting hacked

### Why is a cluster so rewarding to hack

- You have lots and lots of computational power, just be hacking one cluster
- This means, you can even earn money: cryptominer can be installed

### Starting with Tesla

- <https://arstechnica.com/information-technology/2018/02/tesla-cloud-resources-are-hacked-to-run-cryptocurrency-mining-malware/>

## Category 1 by Layer: OS / Kernel

## Securing the OS and the Kernel

### Kernel

- Always patch to the newest kernel
- Be sure to restart the server (in most cases new kernel will start to get used after reboot)

### Good tool for detecting great hardening kernel parameter

- Kernel hardening checker

### Modules

```
## sysctl -w kernel.modules_disabled=1
kernel.modules_disabled = 1
```

```
* If possible harden your kernel, e.g.
* But of course, it is then not allowed to load modules after that isset
```

### OS

- Only install the really needed software in os
  - Eventuall start from a minimal image
- Close unneeded ports

### OS-Patching

- Patch frequently. Eventually using unattended-upgrades

### Hardening Guide

- A bit older, but has really good hints

[Telekom Hardening Guide](#)

### Kernel Hardening Checker

#### Checker

```
https://github.com/a13xp0p0v/kernel-hardening-checker?tab=readme-ov-file
```

```
## Installation
cd /usr/src
git clone https://github.com/a13xp0p0v/kernel-hardening-checker?tab=readme-ov-file
cd kernel-hardening-checker
```

```
sudo sysctl -a > sysctl.file && ./bin/kernel-hardening-checker -c /boot/config-6.8.0-44-generic -l /proc/cmdline -s sysctl.file
```

### Kernel Defence Map

- <https://github.com/a13xp0p0v/linux-kernel-defence-map>

### Guidelines

- <https://gist.github.com/dante-robinson/3a2178e43009c8267ac02387633ff8ca>

## Category 2 by Layer: Cluster

### Securing the components

#### Securing kubelet

#### Breach 1: bypass admission controller

```
If a static Pod fails admission control, the kubelet won't register the
Pod with the API server. However, the Pod still runs on the node.
```

- <https://kubernetes.io/docs/concepts/security/api-server-bypass-risks/>

#### Mitigate breach: Disable the directory for static pods on worker-nodes

- <https://kubernetes.io/docs/concepts/security/api-server-bypass-risks/#static-pods-mitigations>

```
## change the setting kubelet-config
staticPodPath: /etc/kubernetes/manifests
## -> to
staticPodPath:
```

```
after that:
```

- Log in to a kubeadm node
- Run `kubeadm upgrade node phase kubelet-config` to download the latest `kubelet-config` ConfigMap contents into the local file `/var/lib/kubelet/config.yaml`
- Edit the file `/var/lib/kubelet/kubeadm-flags.env` to apply additional configuration with flags
- Restart the kubelet service with `systemctl restart kubelet`

<https://kubernetes.io/docs/tasks/administer-cluster/kubeadm/kubeadm-reconfigure/>

- Only Enable the behaviour really needed
- <https://kubernetes.io/docs/tasks/configure-pod-container/static-pod/#static-pod-creation>

## Breach 2: Allowing anonymous access to kubelet

### Disabling it in anycase

- in the newer of kubeadm it is already the case
- Check with provider / installer

```
authentication:
  anonymous:
    enabled: false
```

- Log in to a kubeadm node
- Run `kubeadm upgrade node phase kubelet-config` to download the latest `kubelet-config` ConfigMap contents into the local file `/var/lib/kubelet/config.yaml`
- Edit the file `/var/lib/kubelet/kubeadm-flags.env` to apply additional configuration with flags
- Restart the kubelet service with `systemctl restart kubelet`


## Least Privileges with RBAC

### The least privileges principles

- Always design your pods, user and components, that they really only have the minimal principles they need
- RBAC Resources help you to do that (Service Accounts, Roles, ClusterRoles, Rolebinding, Clusterrolebindings, Groups)

## Admission Controller

### What does it do ? (The picture)



### How do the docs describe it ?

An admission controller is a piece of code that intercepts requests to the Kubernetes API server prior to persistence of the object, but after the request is authenticated and authorized.

- intercepts request = gets all the requests and validates or changes (=mutates them)
- Reference: <https://kubernetes.io/docs/reference/access-authn-authz/admission-controllers/>

### What kind of admissionControllers do we have ?

- Mutating and Validating
- There are 2 phases of the AdmissionControlProcess: First mutating, then validating

## The static admissionPlugins

- There are static admissionPlugins which are activated by config
- You can see the activated like so

```
## in our system like so.  
kubectl -n kube-system describe pods kube-apiserver-controlplane | grep enable-adm
```

```
--enable-admission-plugins=NodeRestriction
```

- There are some that are activated by default:

```
## in Kubernetes 1.31  
CertificateApproval,  
CertificateSigning,  
CertificateSubjectRestriction,  
DefaultIngressClass,  
DefaultStorageClass,  
DefaultTolerationSeconds,  
LimitRanger,  
MutatingAdmissionWebhook,  
NamespaceLifecycle,  
PersistentVolumeClaimResize,  
PodSecurity,  
Priority,  
ResourceQuota,  
RuntimeClass,  
ServiceAccount,  
StorageObjectInUseProtection,  
TaintNodesByCondition,  
ValidatingAdmissionPolicy,  
ValidatingAdmissionWebhook
```

- Reference: What does which AdmissionController (= AdmissionPlugin) do ? <https://kubernetes.io/docs/reference/access-authn-authz/admission-controllers/#what-does-each-admission-controller-do>

## Category 3 by Layer: Pods Container

### The runAs Options in SecurityContext

- Best practices. Set on level of the pod
- This also reflects containers that are started with kubectl debug (ephemeral containers)

### runAsUser

- Important: UID does not need to exist in container
- Really run as specific user.
- If not set the user from Dockerfile is taken
- Recommended to set it, that will be deep defense line
  - If image has a root user or can not run as root this will fail

### Exercise 1

```
cd  
mkdir -p manifests  
cd manifests  
mkdir run1  
cd run1
```

```
nano 01-pod.yaml
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  creationTimestamp: null  
  labels:  
    run: pod  
  name: nginxrun  
spec:  
  securityContext:  
    runAsUser: 10001  
  containers:  
  - image: nginx:1.23  
    name: pod  
    resources: {}  
  dnsPolicy: ClusterFirst  
  restartPolicy: Always
```



```
kubectl apply -f 01-pod.yaml
kubectl describe pod nginxrun
kubectl logs nginxrun
```

## Exercise 2: (works)

```
cd
mkdir -p manifests
cd manifests
mkdir run2
cd run2
```

```
nano 01-pod.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: alpinerun
spec:
  securityContext:
    runAsUser: 10001
  containers:
  - image: alpine
    command:
      - sleep
      - infinity
    name: pod
```

```
kubectl apply -f .
kubectl describe pod alpinerun
kubectl exec -it alpinerun -- sh
```

```
id
cd /proc/1/ns
ls -la
```

## runAsGroup

- Recommended to set this as well

## runAsNonRoot

- Indicates that use must run as none root
- If this is not configure in the image, the start fails

## sysctls in pods/containers

### What ? Set kernel parameters

- set kernel params in container
- This will only be set for the namespace

### Two groups

- sysctl's that are considered safe
- sysctl's that are considered unsecure (these are not enabled by default in Kubernetes)

## Safe Settings

```
kernel.shm_rmid_forced;
net.ipv4.ip_local_port_range;
net.ipv4.tcp_syncookies;
net.ipv4.ping_group_range (since Kubernetes 1.18);
net.ipv4.ip_unprivileged_port_start (since Kubernetes 1.22);
net.ipv4.ip_local_reserved_ports (since Kubernetes 1.27, needs kernel 3.16+);
net.ipv4.tcp_keepalive_time (since Kubernetes 1.29, needs kernel 4.5+);
net.ipv4.tcp_fin_timeout (since Kubernetes 1.29, needs kernel 4.6+);
net.ipv4.tcp_keepalive_intvl (since Kubernetes 1.29, needs kernel 4.5+);
net.ipv4.tcp_keepalive_probes (since Kubernetes 1.29, needs kernel
```

## Exceptions form Safe Settings

The net.\* sysctls are not allowed with host networking enabled.  
The net.ipv4.tcp\_syncookies sysctl is not namespaced on Linux kernel version 4.5 or lower.

## Example

```

apiVersion: v1
kind: Pod
metadata:
  name: sysctl-example
spec:
  securityContext:
    sysctls:
      - name: kernel.shm_rmid_forced
        value: "0"
      - name: net.core.somaxconn
        value: "1024"
      - name: kernel.msgmax
        value: "65536"

```

## Overview capabilities

### What are these ?

- Capabilities allow us to execute stuff, that normally only the root user can do.

### Best practice for security and container/pods

- Tear capabilities down to Drop: all and set up those, that you need

### As little capabilities as possible.

- Use as little capabilities as possible in your pod/cont

### List of capabilities

- cap\_chown
- cap\_dac\_override
- cap\_fowner
- cap\_fsetid
- cap\_kill
- cap\_setgid
- cap\_setuid
- cap\_setpcap
- cap\_net\_bind\_service
- cap\_net\_raw
- cap\_sys\_chroot
- cap\_mknod
- cap\_audit\_write
- cap\_setfcap

#### cap\_chown

- User can chown (Change Owner without being root)

#### cap\_dac\_override

- Bypasses permission checks

#### cap\_fowner

```

Bypass permission checks on operations that normally
require the filesystem UID of the process to match the
UID of the file (e.g., chmod(2), utime(2)),

```

#### cap\_fsetid

- Safe to disable
- Not needed as we should not use setgid and setuid bits anyway

```

when creating a new folder in folder with setgid, new folder
will have this permission as well

```

#### cap\_kill

- Please not have this enabled.
- User is allowed to kill processes within the container

```

Bypass permission checks for sending signals
## should not be the job of the container

```

#### cap\_setgid

- Do not enable !

```

Allow to change GID of a process

```

#### cap\_setuid

- makes it possible to privilege escalations

```
make arbitrary manipulations of process UID
(setuid(2), setreuid(2), setresuid(2), setfsuid(2));
```

#### cap\_setpcap

Allow user to set other process capabilities

#### cap\_net\_bind\_service

- If you want to bind a privilege port, you will need this
- Like starting httpd on port 80

Security Question/Hint:  
Is there probably a better way to do this:  
e.g. Open Port 8080 -> and having the service on port 80

#### cap\_net\_raw

- Needed to open a raw socket
- Needed to perform ping
- There have been many vulnerabilities concerning this capability in the past e.g. [CVE-2020-14385](#)
- Warning: Do not activate it, if you really, really need this
  - You can ping with a debug container if you need to

#### cap\_sys\_chroot

CAP\_SYS\_CHROOT permits the use of the chroot(2) system call. This may allow escaping of any chroot(2) environment, using known weaknesses and

#### cap\_mknod

- No need to have this
- Allow to create special files under /dev

#### cap\_audit\_write

- Allow to write to kernel audit log

#### cap\_setfcap

- Allow user to set other file capabilities

#### Documentation

- <https://man7.org/linux/man-pages/man7/capabilities.7.html>

#### Start pod without capabilities & how can we see this

##### Exercise 1:

```
cd
mkdir -p manifests
cd manifests
mkdir -p nocap
cd nocap
```

```
nano 01-pod.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nocap-nginx
spec:
  containers:
    - name: web
      image: bitnami/nginx
      securityContext:
        capabilities:
          drop:
            - all
```

```
kubectl apply -f .
kubectl get pods
kubectl logs nocap-nginx
```

##### Exercise 2

```
nano 02-alpine.yaml
```

```

apiVersion: v1
kind: Pod
metadata:
  name: nocap-alpine
spec:
  containers:
    - name: web
      command:
        - sleep
        - infinity
      image: alpine
      securityContext:
        capabilities:
          drop:
            - all

```

```

kubectl apply -f .
kubectl get pods
kubectl logs nocap-alpine
kubectl exec -it nocap-alpine -- sh

```

```

ping www.google.de
wget -O - http://www.google.de

```

## Hacking and exploration session HostPID

### Great but still alpha User Namespaces

## Reaction

### The Audit Logs

#### Hints (this is on a system, where kubernetes-api-server runs as static pod)

- When the config unter /etc/kubernetes/manifests/kube-apiserver.yaml changes
  - kubelet automatically detects this and restarts the server
  - if there is a misconfig the pod will vanish

```
## There are 4 stages, that can be monitored:
```

- **RequestReceived** - The stage for events generated as soon as the audit handler receives the request, and before it is delegated down the handler chain.
- **ResponseStarted** - Once the response headers are sent, but before the response body is sent. This stage is only generated for long-running requests (e.g. watch).
- **ResponseComplete** - The response body has been completed and no more bytes will be sent.
- **Panic** - Events generated when a panic occurred.

#### Step 1: 1st -> session (on control plane): Watch kube-apiserver - pod on controlplane

```
## we want to
watch crictl pods | grep api

```

#### Step 2: 2nd -> session (on control plane): create a policy

```
cd /etc/kubernetes
```

```
nano audit-policy.yaml
```

```

apiVersion: audit.k8s.io/v1 # This is required.
kind: Policy
## Don't generate audit events for all requests in RequestReceived stage.
omitStages:

```

```

- "RequestReceived"
rules:
  # Log pod changes at RequestResponse level
  - level: RequestResponse
    resources:
      - group: ""
        # Resource "pods" doesn't match requests to any subresource of pods,
        # which is consistent with the RBAC policy.
        resources: ["pods"]
  # Log "pods/log", "pods/status" at Metadata level
  - level: Metadata
    resources:
      - group: ""
        resources: ["pods/log", "pods/status"]

  # Don't log requests to a configmap called "controller-leader"
  - level: None
    resources:
      - group: ""
        resources: ["configmaps"]
        resourceNames: ["controller-leader"]

  # Don't log watch requests by the "system:kube-proxy" on endpoints or services
  - level: None
    users: ["system:kube-proxy"]
    verbs: ["watch"]
    resources:
      - group: "" # core API group
        resources: ["endpoints", "services"]

  # Don't log authenticated requests to certain non-resource URL paths.
  - level: None
    userGroups: ["system:authenticated"]
    nonResourceURLs:
      - "/api*" # Wildcard matching.
      - "/version"

  # Log the request body of configmap changes in kube-system.
  - level: Request
    resources:
      - group: "" # core API group
        resources: ["configmaps"]
    # This rule only applies to resources in the "kube-system" namespace.
    # The empty string "" can be used to select non-namespaced resources.
    namespaces: ["kube-system"]

  # Log configmap and secret changes in all other namespaces at the Metadata level.
  - level: Metadata
    resources:
      - group: "" # core API group
        resources: ["secrets", "configmaps"]

  # Log all other resources in core and extensions at the Request level.
  - level: Request
    resources:
      - group: "" # core API group
      - group: "extensions" # Version of group should NOT be included.

  # A catch-all rule to log all other requests at the Metadata level.
  - level: Metadata
    # Long-running requests like watches that fall under this rule will not
    # generate an audit event in RequestReceived.
    omitStages:
      - "RequestReceived"

```

```

## Important: You do not need to apply/create it.

```

### Step 3: 2nd -> session: Change settings in /etc/kubernetes/manifests/kube-apiserver.yaml

```

## security copy
cp /etc/kubernetes/manifests/kube-apiserver.yaml /root

## Add lines in /etc/kubernetes/manifests/kube-apiserver.yaml
- --audit-log-path=/var/log/kubernetes/apiserver/audit/audit.log
- --audit-policy-file=/etc/kubernetes/audit-policies.yaml

```

```
## Add lines under volumeMounts
- mountPath: /etc/kubernetes/audit-policy.yaml
  name: audit
  readOnly: true
- mountPath: /var/log/kubernetes/apiserver/audit/
  name: audit-log
  readOnly: false
```

```
## Add volumes lines under volumes
- name: audit
  hostPath:
    path: /etc/kubernetes/audit-policy.yaml
    type: File

- name: audit-log
  hostPath:
    path: /var/log/kubernetes/audit/
    type: DirectoryOrCreate
```

#### Step 4: 1st -> session

POD ID	CREATED	STATE	NAME	NAMESPACE	A
27bfc9f8d1552	RUNTIME 7 seconds ago (default)	Ready	kube-apiserver-controlplane	kube-system	0
43e9ba82707c0	2 hours ago	Ready	csi-node-driver-28l2j	calico-system	1

c

#### Step 5: 2nd -> session

```
## There should be enough noise already
cat /var/log/kubernetes/audit/audit.log
```

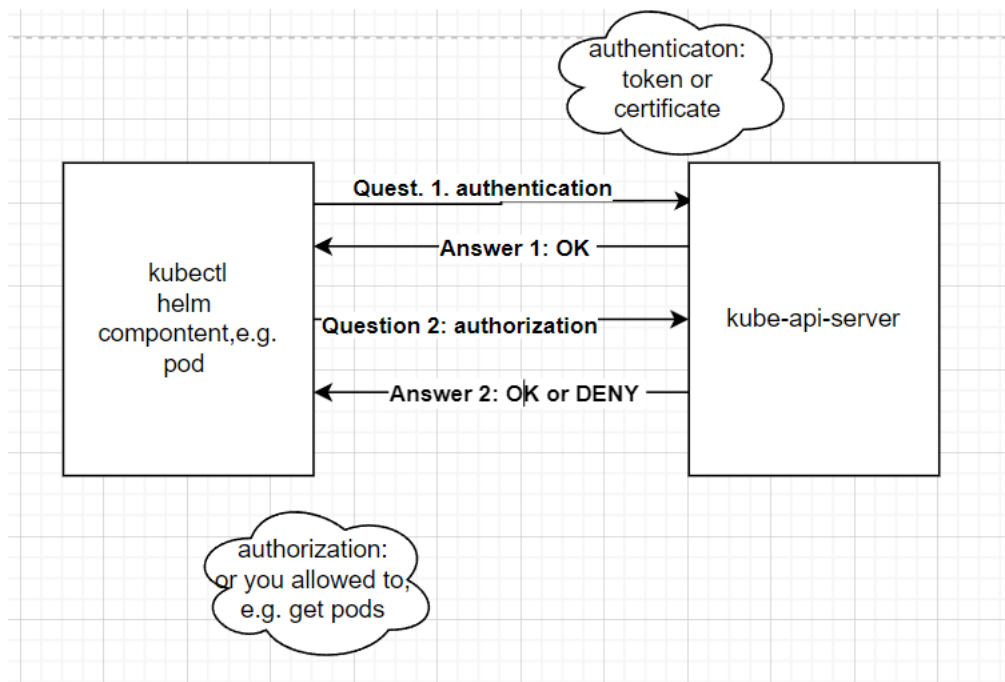
#### Reference

- <https://kubernetes.io/docs/tasks/debug/debug-cluster/audit/>

## RBAC

#### How does RBAC work ?

- Let us see in a picture



#### Where does RBAC play a role ?

Users -> kube-api-server

- User how want to access the kube api server

## Components -> kube-api-server

- e.g. kubelet -> kube-api-server

## Pods / System Pods -> kube-api-server

- Pods and System Pods (e.g. kube-proxy a.k.a. CoreDNS) how want to access the kube-api-server

## kubeconfig decode certificate

```
cd
cd .kube
cat config
## copy client-certificate-data

echo
LS0tLS1CRUdJTiBDRVJUSUZJQ0FURSB0tLS0tCk1JSURLVENDQWhHZ0F3SUJBZ0lJVGpPWWowbXpWME13RFFZSktvWk1odmNOQVFFTEJRQXdGVEVUTUJFR0ExVUUKQXhNS2Ez
base64 -d > out.crt
openssl x509 -in out.crt -text -noout
```

## kubectl check your permission - can-i

### A specific command

```
kubectl auth can-i get pods
```

### List all

```
kubectl auth can-i --list
```

## use kubectl in pod - default service account

### Walkthrough

```
kubectl run -it --rm kubectltester --image=alpine -- sh
```

```
## in shell
apk add kubectl
## it uses in in-cluster configuration in folder
## /var/run/secrets/kubernetes.io/serviceaccount
kubectl auth can-i --list
```

## create user for kubeconfig with using certificate

### Step 0: create an new rolebinding for the group (we want to use)

```
kubectl create rolebinding developers --clusterrole=view --group=developers
```

### Step 1: on your client: create private certificate

```
cd
mkdir -p certs
## create your private key
openssl genrsa -out ~/certs/jochen.key 4096
```

### Step 2: on your client: create csr (certificate signing request)

```
nano ~/certs/jochen.csr.cnf
```

```
[ req ]
default_bits = 2048
prompt = no
default_md = sha256
distinguished_name = dn
[ dn ]
CN = jochen
O = developers
[ v3_ext ]
authorityKeyIdentifier=keyid,issuer:always
basicConstraints=CA:FALSE
keyUsage=keyEncipherment,dataEncipherment
extendedKeyUsage=serverAuth,clientAuth
```

```
## Create Certificate Signing Request
openssl req -config ~/certs/jochen.csr.cnf -new -key ~/certs/jochen.key -nodes -out ~/certs/jochen.csr
openssl req -in certs/jochen.csr --noout -text
```

### Step 3: Send approval request to server

```
## get csr (base64 decoded)
cat ~/certs/jochen.csr | base64 | tr -d '\n'
```

```
cd certs
nano jochen-csr.yaml
```

```
apiVersion: certificates.k8s.io/v1
kind: CertificateSigningRequest
metadata:
  name: jochen-authentication
spec:
  signerName: kubernetes.io/kube-apiserver-client
  groups:
    - system:authenticated
  request:
LS0tLS1CRUdJTiBDRVJUSUZJQ0FURSBRSRVFVRNULS0tLS0KTUJRWf6Q0NbBE1DQVFBd0pqRVBNQTBHQTfVRUF3d0d0hbTlqYUdwdU1STXdfUVlEVlFRS0RBcGtaWfPsYkc5

  usages:
    - client auth
```

```
kubectl apply -f jochen-csr.yaml
kubectl get -f jochen-csr.yaml
## show me the current state -> pending
kubectl describe -f jochen-csr.yaml
```

### Step 4: approve signing request

```
kubectl certificate approve jochen-authentication
## or:
kubectl certificate approve -f jochen-csr.yaml

## see, that it is approved
kubectl describe -f jochen-csr.yaml
```

### Step 5: get the approved certificate to be used

```
kubectl get csr jochen-authentication -o jsonpath='{.status.certificate}' | base64 --decode > ~/certs/jochen.crt
```

### Step 6: construct kubeconfig for new user

```
cd
cd certs
```

```
## create new user
kubectl config set-credentials jochen --client-certificate=jochen.crt --client-key=jochen.key
```

```
## add a new context
kubectl config set-context jochen --user=jochen --cluster=kubernetes
```

### Step 7: Use and test the new context

```
kubectl config use-context jochen
kubectl get pods
```

#### Ref:

- <https://kb.leaseweb.com/kb/users-roles-and-permissions-on-kubernetes-rbac/kubernetes-users-roles-and-permissions-on-kubernetes-rbac-create-a-certificate-based-kubeconfig/>

### practical exercise rbac

#### Schritt 1: Create a service account and a secret

```
cd
mkdir -p manifests/rbac
cd manifests/rbac
```

#### Mini-Step 1: definition of the user



```
## vi service-account.yml
apiVersion: v1
kind: ServiceAccount
metadata:
  name: training
  namespace: default
```

```
kubectl apply -f service-account.yml
```

#### Mini-Step 1.5: create Secret

- From Kubernetes 1.25 tokens are not created automatically when creating a service account (sa)
- You have to create them manually with annotation attached
- <https://kubernetes.io/docs/reference/access-authn-authz/service-accounts-admin/#create-token>

```
## vi secret.yml
apiVersion: v1
kind: Secret
type: kubernetes.io/service-account-token
metadata:
  name: trainingtoken
  annotations:
    kubernetes.io/service-account.name: training
```

```
kubectl apply -f .
```

#### Mini-Schritt 2: ClusterRole creation - Valid for all namespaces but it has to get assigned to a clusterrolebinding or rolebinding

```
### Does not work, before there is no assignment
## vi pods-clusterrole.yml
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: pods-clusterrole
rules:
- apiGroups: [""] # "" indicates the core API group
  resources: ["pods"]
  verbs: ["get", "watch", "list", "create"]
```

```
kubectl apply -f pods-clusterrole.yml
```

#### Mini-Schritt 3: Assigning the clusterrole to a specific service account

```
## vi rb-training-ns-default-pods.yml
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: rolebinding-ns-default-pods
  namespace: default
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: pods-clusterrole
subjects:
- kind: ServiceAccount
  name: training
  namespace: default
```

```
kubectl apply -f rb-training-ns-default-pods.yml
```

#### Mini-Step 4: Test it (does the access work)

```
kubectl auth can-i get pods -n default --as system:serviceaccount:default:training
kubectl auth can-i --list --as system:serviceaccount:default:training
```

#### Schritt 2: create Context / read Credentials and put them in kubeconfig (the Kubernetes-Version 1.25. way)

##### Mini-Step 1: kubeconfig setzen

```
kubectl config set-context training-ctx --cluster kubernetes --user training

## extract name of the token from here

TOKEN=`kubectl get secret trainingtoken -o jsonpath='{.data.token}' | base64 --decode`
echo $TOKEN
kubectl config set-credentials training --token=$TOKEN
```

```
kubectl config use-context training-ctx

## Hier reichen die Rechte nicht aus
kubectl get deploy
## Error from server (Forbidden): pods is forbidden: User "system:serviceaccount:kube-system:training" cannot list # resource
"pods" in API group "" in the namespace "default"
```

#### Mini-Step 2:

```
kubectl config use-context training-ctx
kubectl get pods
```

#### Mini-Step 3: back to the old context

```
kubectl config get-contexts

kubectl config use-context cluster-admin@kubernetes
```

#### Refs:

- <https://docs.oracle.com/en-us/iaas/Content/ContEng/Tasks/contengaddingerviceacctoken.htm>
- <https://microk8s.io/docs/multi-user>
- <https://faun.pub/kubernetes-rbac-use-one-role-in-multiple-namespaces-d1d08bb08286>


#### Ref: Create Service Account Token

- <https://kubernetes.io/docs/reference/access-authn-authz/service-accounts-admin/#create-token>

## Obey Security Policies (AdmissionControllers)

### Admission Controller

#### What does it do ? (The picture)



#### How do the docs describe it ?

An admission controller is a piece of code that intercepts requests to the Kubernetes API server prior to persistence of the object, but after the request is authenticated and authorized.

- intercepts request = gets all the requests and validates or changes (=mutates them)
- Reference: <https://kubernetes.io/docs/reference/access-authn-authz/admission-controllers/>

#### What kind of admissionControllers do we have ?

- Mutating and Validating
- There are 2 phases of the AdmissionControlProcess: First mutating, then validating

#### The static admissionPlugins

- There are static admissionPlugins which are activated by config
- You can see the activated like so

```
## in our system like so.
kubectl -n kube-system describe pods kube-apiserver-controlplane | grep enable-adm
```

```
--enable-admission-plugins=NodeRestriction
```

- There are some that are activated by default:

```
## in Kubernetes 1.31
CertificateApproval,
CertificateSigning,
CertificateSubjectRestriction,
```

```

DefaultIngressClass,
DefaultStorageClass,
DefaultTolerationSeconds,
LimitRanger,
MutatingAdmissionWebhook,
NamespaceLifecycle,
PersistentVolumeClaimResize,
PodSecurity,
Priority,
ResourceQuota,
RuntimeClass,
ServiceAccount,
StorageObjectInUseProtection,
TaintNodesByCondition,
ValidatingAdmissionPolicy,
ValidatingAdmissionWebhook

```

- Reference: What does which AdmissionController (= AdmissionPlugin) do ? <https://kubernetes.io/docs/reference/access-authn-authz/admission-controllers/#what-does-each-admission-controller-do>

### Exercise with PSA

#### Seit: 1.2.22 Pod Security Admission

- 1.2.22 - Alpha Feature, was not activated by default. need to activate it as feature gate (Kind)
- 1.2.23 - Beta -> probably

#### Predefined settings

- privileges - no restrictions
- baseline - some restriction
- restricted - really restrictive
- Reference: <https://kubernetes.io/docs/concepts/security/pod-security-standards/>

### Practical Example starting from Kubernetes 1.23

```

mkdir -p manifests
cd manifests
mkdir psa
cd psa
nano 01-ns.yml

```

```

## Step 1: create namespace

apiVersion: v1
kind: Namespace
metadata:
  name: test-ns1
  labels:
    # soft version - running but showing complaints
    # pod-security.kubernetes.io/enforce: baseline
    pod-security.kubernetes.io/enforce: restricted
    pod-security.kubernetes.io/audit: restricted
    pod-security.kubernetes.io/warn: restricted

```

```

kubectl apply -f 01-ns.yml

```

```

## Schritt 2: Testen mit nginx - pod
## vi 02-nginx.yml

```

```

apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: test-ns1
spec:
  containers:
    - image: nginx
      name: nginx
      ports:
        - containerPort: 80

```

```

## a lot of warnings will come up
## because this image runs as root !! (by default)
kubectl apply -f 02-nginx.yml

```

```
## Schritt 3:  
## Change SecurityContext in Container
```

```
## vi 02-nginx.yml  
  
apiVersion: v1  
kind: Pod  
metadata:  
  name: nginx  
  namespace: test-ns1  
spec:  
  containers:  
  - image: nginx  
    name: nginx  
    ports:  
    - containerPort: 80  
    securityContext:  
      seccompProfile:  
        type: RuntimeDefault
```

```
kubectl delete -f 02-nginx.yml  
kubectl apply -f 02-nginx.yml  
kubectl -n test-ns1 get pods
```

```
## Schritt 4:  
## Weitere Anpassung runAsNotRoot  
## vi 02-nginx.yml  
apiVersion: v1  
kind: Pod  
metadata:  
  name: nginx  
  namespace: test-ns<tln>  
spec:  
  containers:  
  - image: nginx  
    name: nginx  
    ports:  
    - containerPort: 80  
    securityContext:  
      seccompProfile:  
        type: RuntimeDefault  
      runAsNonRoot: true
```

```
## pod kann erstellt werden, wird aber nicht gestartet  
kubectl delete -f 02-nginx.yml  
kubectl apply -f 02-nginx.yml  
kubectl -n test-ns1 get pods  
kubectl -n test-ns1 describe pods nginx
```

```
## Schritt 4:  
## Anpassen der Sicherheitseinstellung (Phase1) im Container
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: nginx  
  namespace: test-ns1  
spec:  
  containers:  
  - image: nginx  
    name: nginx  
    ports:  
    - containerPort: 80  
    securityContext:  
      seccompProfile:  
        type: RuntimeDefault  
      runAsNonRoot: true  
      allowPrivilegeEscalation: false  
      capabilities:  
        drop: ["ALL"]
```

```
kubectl delete -f 02-nginx.yml  
kubectl apply -f 02-nginx.yml  
kubectl -n test-ns1 get pods
```

## Praktisches Beispiel für Version ab 1.2.23 -Lösung - Container als NICHT-Root laufen lassen

- Wir müssen ein image, dass auch als NICHT-Root laufen kann
- ..oder selbst eines bauen (:o)) o bei nginx ist das bitnami/nginx

```
## vi 03-nginx-bitnami.yml
apiVersion: v1
kind: Pod
metadata:
  name: bitnami-nginx
  namespace: test-ns1
spec:
  containers:
    - image: bitnami/nginx
      name: bitnami-nginx
      ports:
        - containerPort: 80
      securityContext:
        seccompProfile:
          type: RuntimeDefault
        runAsNonRoot: true
```

```
## und er läuft als nicht root
kubectl apply -f 03_pod-bitnami.yml
kubectl -n test-ns1 get pods
```

## OPA Gatekeeper 01-Overview

### How does it work ?

- It is called by the definition in mutationAdmissionWebhook and validatingAdmissionWebhook

### What can the OPA Gatekeeper do ?

- It can validate
- It can mutate

### How does it do this ?

- It uses a language which is called REGO
- It uses objects like

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate

apiVersion: templates.gatekeeper.sh/v1
kind: Constraint
```

- for the validation

## OPA Gatekeeper 02-Install with Helm

### Step 1: Installation (helm)

```
helm repo add gatekeeper https://open-policy-agent.github.io/gatekeeper/charts
helm upgrade gatekeeper/gatekeeper --install gatekeeper --namespace gatekeeper-system --create-namespace
```

### Step 2: Webhooks (lookaround)

- This create a mutation and a validationWebhook

```
kubectl get validatingwebhookconfigurations gatekeeper-validating-webhook-configuration
kubectl get mutatingwebhookconfigurations gatekeeper-mutating-webhook-configuration
```

- Let's look in the mutation more deeply

```
kubectl get mutatingwebhookconfigurations gatekeeper-mutating-webhook-configuration -o yaml
```

### Step 3: The components

```
## controllers are the endpoint for the webhooking
## audit is done every 60 seconds in the audit-pod
kubectl -n gatekeeper-system get all
```

## OPA Gatekeeper 03-Simple Exercise

### Step 1: Create constraintTemplate

- I took this from the library: <https://open-policy-agent.github.io/gatekeeper-library/website/>

- <https://open-policy-agent.github.io/gatekeeper-library/website/validation/block-nodeport-services/>

```
cd
mkdir -p manifests
cd manifests
mkdir restrict-node-port
cd restrict-node-port
```

```
nano 01-constraint-template.yaml
```

```
apiVersion: templates.gatekeeper.sh/v1
kind: ConstraintTemplate
metadata:
  name: k8sblocknodeport
  annotations:
    metadata.gatekeeper.sh/title: "Block NodePort"
    metadata.gatekeeper.sh/version: 1.0.0
  description: >-
    Disallows all Services with type NodePort.

    https://kubernetes.io/docs/concepts/services-networking/service/#nodeport
spec:
  crd:
    spec:
      names:
        kind: K8sBlockNodePort
      targets:
        - target: admission.k8s.gatekeeper.sh
          rego: |
            package k8sblocknodeport

            violation[{"msg": msg}] {
              input.review.kind.kind == "Service"
              input.review.object.spec.type == "NodePort"
              msg := "User is not allowed to create service of type NodePort"
            }
```

```
kubectl apply -f .
```

## Step 2: Create constraint

- it is like an instance (in code = usage of classes, can be created multiple times)
- the match defines, when it triggers -> when it calls the constraintTemplate for validation

```
nano 02-constraint.yaml
```

```
apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sBlockNodePort
metadata:
  name: block-node-port
spec:
  match:
    kinds:
      - apiGroups: [""]
        kinds: ["Service"]
```

```
kubectl apply -f .
```

## Step 3: Test constraint with Service

```
nano 03-service.yaml
```

```
apiVersion: v1
kind: Service
metadata:
  name: my-service-disallowed
spec:
  type: NodePort
  ports:
    - port: 80
```

```
kubectl apply -f .
```

```
Error from server (Forbidden): error when creating "03-service.yaml":
admission webhook "validation.gatekeeper.sh" denied the request:
```

```
[block-node-port] User is not allowed to create service of type NodePort
```

## OPA Gatekeeper 04-Example-Job-Debug

### Step 1: Create constraintTemplate

```
cd
mkdir -p manifests
cd manifests
mkdir blockjob
cd blockjob
```

```
nano constraint-template.yaml
```

```
kind: ConstraintTemplate
metadata:
  name: k8sblockjob
  annotations:
    metadata.gatekeeper.sh/title: "Block Job"
    metadata.gatekeeper.sh/version: 1.0.0
  description: >-
    Blocks certain jobs
spec:
  crd:
    spec:
      names:
        kind: K8sBlockJob
      targets:
        - target: admission.k8s.gatekeeper.sh
          rego: |
            package k8sblockjob

            violation[{"msg": msg}] {
              # input.review.kind.kind == "Job"
              msg1:= sprintf("Data: %v", [input.review.userInfo])
              msg2:= sprintf("JOBS not allowed .. REVIEW OBJECT: %v", [input.review])

              msg:= concat(" ", [msg1,msg2])
            }
```

```
kubectl apply -f .
## Was it successfully parsed and compiled ?
kubectl describe -f constraint-template.yaml
```

### Step 2: Create constraint

```
nano constraint.yaml
```

```
cat 02-constraint.yaml
apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sBlockJob
metadata:
  name: block-job
spec:
  match:
    kinds:
      # Important batch not Batch, Batch will not work
      - apiGroups: ["batch"]
        kinds: ["Job"]
```

```
kubectl apply -f .
```

### Step 3: Test with Job

```
nano 03-job.yaml
```

```
apiVersion: batch/v1
kind: Job
metadata:
  name: pi
spec:
  template:
    spec:
      containers:
        - name: pi
```

```
    image: perl:5.34.0
    command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(2000)"]
    restartPolicy: Never
    backoffLimit: 4
```

Should not work:

```
kubectl apply -f .
```

#### Step 4: Let's try from a pod

- Prepare user

```
kubectl create sa podjob
kubectl create rolebinding podjob-binding --clusterrole=cluster-admin --serviceaccount=default:podjob
kubectl run -it --rm jobmaker --image=alpine --overrides='{ "spec": { "serviceAccount": "podjob" } }' -- sh
```

```
## in pod install kubectl and nano
apk add nano kubectl
## create pod manifests
cd
nano job.yaml
```

```
apiVersion: batch/v1
kind: Job
metadata:
  name: pi
spec:
  template:
    spec:
      containers:
      - name: pi
        image: perl:5.34.0
        command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(2000)"]
        restartPolicy: Never
      backoffLimit: 4
```

```
## should not work
kubectl apply -f pod.yaml
```

#### Connaisseur: Verifying images before Deployment

##### Prerequisites

- You must have create a private/public key pair
- You must have signed some images in your registry on docker Hub.
- This was done here: [Signing an image with cosign](#)

##### Step 1: Install Connaissuer with helm

```
cd
mkdir -p manifests
cd manifests
mkdir connaisseur
cd connaisseur
```

```
nano values.yaml
```

```
## 1. We will add the public key of ours in validators: cosign->type:cosign
## 2. We will add a policy for the system to know, when to use it:
- pattern: "docker.io/dockertrainereu/*:*"
  validator: cosign
## Unfortunately we muss add everything from the defaults values file
## concerning -> validators, policy
## I have tested this ....
```

```
application:
  features:
    namespacedValidation:
      mode: validate

## validator options: https://sse-secure-systems.github.io/connaisseur/latest/validators/
validators:
- name: allow
  type: static
  approve: true
```



```

- name: deny
  type: static
  approve: false
- name: cosign
  type: cosign
  trustRoots:
    - name: default
      key: |
        -----BEGIN PUBLIC KEY-----
        MFkwEwYHkoZIZj0CAQYIKoZIZj0DAQcDQgAE2LmRkI8sNi6fSluqU5UEisptlwZ1
        YaJBAJqTf96ccM4R3MstL8PfR5fhy877TG7bnpc4YnlfejT6F7XE71FWkA==
        -----END PUBLIC KEY-----
- name: dockerhub
  type: notaryv1
  trustRoots:
    - name: default # root from dockerhub
      key: |
        -----BEGIN PUBLIC KEY-----
        MFkwEwYHkoZIZj0CAQYIKoZIZj0DAQcDQgAEOXYta5TgdCwXTCnLU09W5T4M4r9f
        QQrQJuADP6U7g5r9ICgPSmZuRHP/1AYUfQW3baveKsT969EfELKj1l1fCA==
        -----END PUBLIC KEY-----
    - name: sse # root from sse
      key: |
        -----BEGIN PUBLIC KEY-----
        MFkwEwYHkoZIZj0CAQYIKoZIZj0DAQcDQgAesx28WV7BsQfnHF1kZmpdCTLJaWe
        d0CA+Joi8H4REuBaWSZ5zPDe468WuOJ6f71E7WFG3CCEVYHuoZt2UYbN/Q==
        -----END PUBLIC KEY-----

policy:
- pattern: ".*"
  validator: deny
- pattern: "docker.io/dockertrainereu/.*"
  validator: cosign
- pattern: "docker.io/library/.*"
  validator: dockerhub
- pattern: "docker.io/securesystemsengineering/.*"
  validator: dockerhub
  with:
    trustRoot: sse
- pattern: "registry.k8s.io/.*"
  validator: allow

```

```

## Add the repo
helm repo add connaisseur https://sse-secure-systems.github.io/connaisseur/charts
## Install the helm chart
helm upgrade connaisseur connaisseur/connaisseur --install --create-namespace --namespace connaisseur -f values.yaml

```

## Step 2: Create a namespace and label it with connaisseur

- In our example we apply it only in a specific namespace
- But you can also use it for all namespaces
- According to: [https://sse-secure-systems.github.io/connaisseur/latest/features/namespaced\\_validation/](https://sse-secure-systems.github.io/connaisseur/latest/features/namespaced_validation/)

```

## create namespace
kubectl create ns appl
kubectl label ns appl securesystemsengineering.connaisseur/xwebhook=validate

```

## Step 3: Try to run image in namespace

```

## image from docker -> works
kubectl -n appl run nginxme --image=nginx:1.23

## signed image from dockertrainereu
kubectl -n appl run pod1 --image=dockertrainereu/alpine-rootless:1.20

## unsigned image from dockertrainereu
kubectl -n appl run pod1 --image=dockertrainereu/pinger

```

## Pod Security

### Automount ServiceAccounts or not ?

#### Why ?

- Every attacker tries to get as much information as possible
- Although there are not severe permissions in here, show as little information as possible
- For example, use will see, which namespace he is in ;o)

## Disable ?

```
## enabled by default
kubectl explain pod.spec.automountServiceAccountToken
```

## Unprivilegierte Pods/Container

### Which images to use ?

#### docker hub

- bitnami images
- Also search for unprivileged -> e.g. <https://hub.docker.com/search?q=unprivileged>
  - BUT: be careful whom you trust

## The SecurityContext

## Network Policies

### Exercise NetworkPolicies

#### Step 1: Create Deployment and Service

```
SHORT=jm
kubectl create ns policy-demo-$SHORT
```

```
cd
mkdir -p manifests
cd manifests
mkdir -p np
cd np
```

```
nano 01-deployment.yml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 1
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.23
          ports:
            - containerPort: 80
```

```
kubectl -n policy-demo-$SHORT apply -f .
```

```
nano 02-service.yml
```

```
apiVersion: v1
kind: Service
metadata:
  name: nginx
spec:
  type: ClusterIP # Default Wert
  ports:
    - port: 80
      protocol: TCP
  selector:
    app: nginx
```

```
kubectl -n policy-demo-$SHORT apply -f .
```

#### Step 2: Testing access without any rules

```
## Run a 2nd pod to access nginx
kubectl run --namespace=policy-demo-$SHORT access --rm -ti --image busybox
```

```
## Within the shell/after prompt
wget -q nginx -O -
```

```
## Optional: Show pod in second 2. ssh-session on jump-host
kubectl -n policy-demo-$SHORT get pods --show-labels
```

### Step 3: Define policy: no access is allowed by default (in this namespace)

```
nano 03-default-deny.yaml
```

```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: default-deny
spec:
  podSelector:
    matchLabels: {}
```

```
kubectl -n policy-demo-$SHORT apply -f .
```

### Step 4: Test connection with deny all rules

```
kubectl run --namespace=policy-demo-$SHORT access --rm -ti --image busybox
```

```
## Within the shell
wget -q nginx -O -
```

### Step 5: Allow access von pods mit dem Label run=access (alle mit run gestarteten pods mit namen access haben dieses label per default)

```
nano 04-access-nginx.yaml
```

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: access-nginx
spec:
  podSelector:
    matchLabels:
      app: nginx
  ingress:
    - from:
      - podSelector:
          matchLabels:
            run: access
```

```
kubectl -n policy-demo-$SHORT apply -f .
```

### Schritt 5: Test it (access should work)

```
## Start a 2nd pod to access nginx
## because of run->access pod automatically has the label run:access
kubectl run --namespace=policy-demo-$SHORT access --rm -ti --image busybox
```

```
## innerhalb der shell
wget -q nginx -O -
```

### Step 6: start Pod with label run=no-access - this should not work

```
kubectl run --namespace=policy-demo-$SHORT no-access --rm -ti --image busybox
```

```
## in der shell
wget -q nginx -O -
```

### Step 7: Cleanup

```
kubectl delete ns policy-demo-$SHORT
```

### Ref:

- <https://projectcalico.docs.tigera.io/security/tutorials/kubernetes-policy-basic>

## ServiceMesh

### Why a ServiceMesh ?

#### What is a service mesh ?

```
A service mesh is an infrastructure layer
that gives applications capabilities
like zero-trust security, observability,
and advanced traffic management, without code changes.
```

#### Advantages / Features

1. Observability & monitoring
2. Traffic management
3. Resilience & Reliability
4. Security
5. Service Discovery

#### Observability & monitoring

- Service mesh offers:
  - valuable insights into the communication between services
  - effective monitoring to help in troubleshooting application errors.

#### Traffic management

- Service mesh offers:
  - intelligent request distribution
  - load balancing,
  - support for canary deployments.
  - These capabilities enhance resource utilization and enable efficient traffic management

#### Resilience & Reliability

- By handling retries, timeouts, and failures,
  - service mesh contributes to the overall stability and resilience of services
  - reducing the impact of potential disruptions.

#### Security

- Service mesh enforces security policies, and handles authentication, authorization, and encryption
  - ensuring secure communication between services and eventually, strengthening the overall security posture of the application.

#### Service Discovery

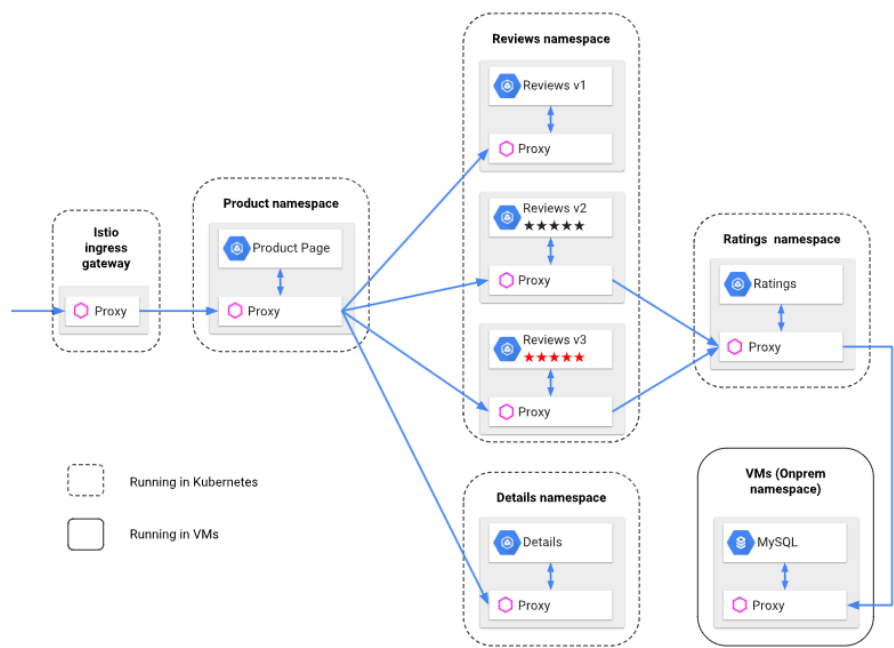
- With service discovery features, service mesh can simplify the process of locating and routing services dynamically
- adapting to system changes seamlessly. This enables easier management and interaction between services.

#### Overall benefits

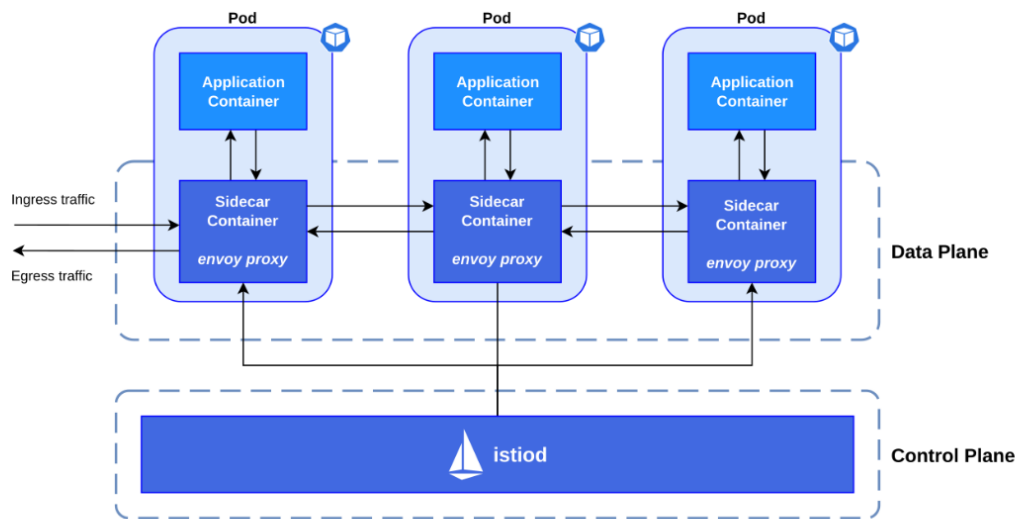
```
Microservices communication:
Adopting a service mesh can simplify the implementation of a microservices architecture by abstracting away infrastructure complexities.
It provides a standardized approach to manage and orchestrate communication within the microservices ecosystem.
```

#### How does a ServiceMeshs work? (example istio

# Overview




## Istio control plane and data plane



• Source: kubebyexample.com

## istio security features

### Overview



Security overview

#### Security needs of microservices

- To defend against man-in-the-middle attacks, they need traffic encryption.
- To provide flexible service access control, they need mutual TLS and fine-grained access policies.
- To determine who did what at what time, they need auditing tools.

#### Implementation of security

The Istio security features provide strong identity, powerful policy, transparent TLS encryption, and authentication, authorization and audit (AAA) tools to protect your services and data. The goals of Istio security are:

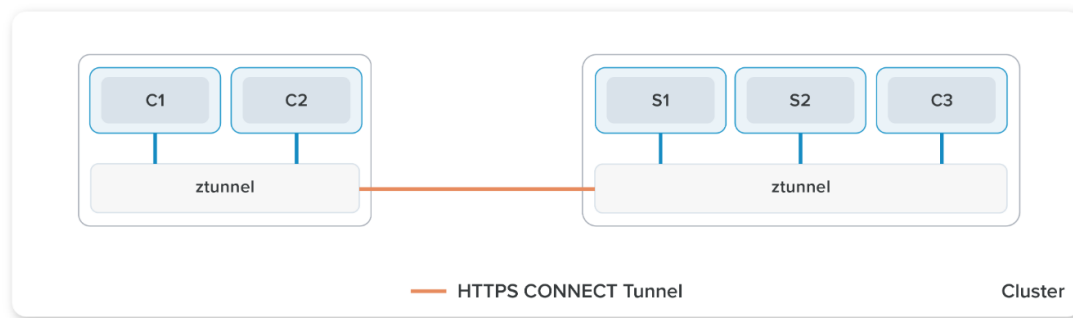
- Security by default: no changes needed to application code and infrastructure
- Defense in depth: integrate with existing security systems to provide multiple layers of defense
- Zero-trust network: build security solutions on distrusted networks

#### istio-service mesh - ambient mode

##### Light: Only Layer 4 per node (ztunnel)

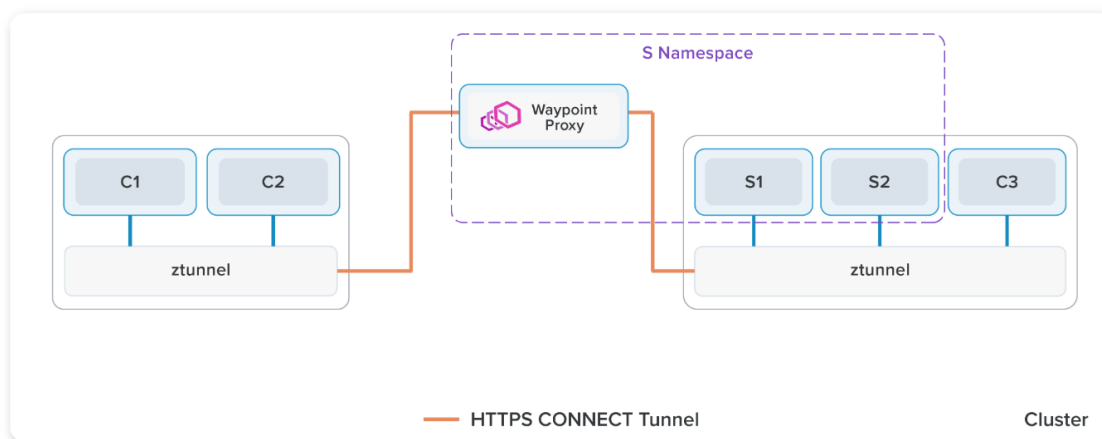
- No sidecar (envoy-proxy) per Pod, but one ztunnel agent per Node (Layer 4)
- Enables security features (mtls, traffic encryption)

Like so:



#### Full fledged: Layer 4 (ztunnel) per Node & Layer 7 per Namespace (

- One waypoint - proxy is rolled out per Namespace, which connects to the ztunnel agents



When additional features are needed, ambient mesh deploys waypoint proxies, which ztunnels connect through for policy enforcement

#### Features in "fully-fledged" - ambient - mode

Application deployment use case	Ambient mode configuration
Zero Trust networking via mutual-TLS, encrypted and tunneled data transport of client application traffic, L4 authorization, L4 telemetry	ztunnel only (default)
As above, plus advanced Istio traffic management features (including L7 authorization, telemetry and VirtualService routing)	ztunnel and waypoint proxies

#### Advantages:

- Less overhead
- One can start step-by-step moving towards a mesh (Layer 4 firstly and if wanted Layer 7 for specific namespaces)
- Old pattern: sidecar and new pattern: ambient can live side by side.

#### Performance comparison - baseline,sidecar,ambient

- <https://liveswayer.io/blog/2024/06/06/comparison-of-service-meshes-part-two/>

#### Image Security

#### Documentation

Great video about attacking kubernetes - older, but some stuff is still applicable

- <https://www.youtube.com/watch?v=HmoVSmTIOxM>

#### **Straight forward hacking session of kubernetes**

- [https://youtu.be/iD\\_klswHJQs?si=97rWNUAbGjLwCjpa](https://youtu.be/iD_klswHJQs?si=97rWNUAbGjLwCjpa)

#### **github with manifests for creating bad pods**

- <https://bishopfox.com/blog/kubernetes-pod-privilege-escalation#pod8>